

Credit Scoring Dataset

The dataset being used describes the customers (seniority, age, marital status, income, and other characteristics), the loan (the requested amount, the price of the idem) as well as its status (was it paid back or not).

The Dataset is available on GitHub at <https://github.com/gastonstat/CreditScoring/>
(<https://github.com/gastonstat/CreditScoring/>)

This is a binary classification problem, and three kinds of algorithms/tools will be utilized:

1. DecisionTreeClassifier (from scikit-learn)
2. RandomForestClassifier (from scikit-learn)
3. Extreme Gradient Boosting (from xgboost library)

```
In [78]: import os, sys, io, requests
import numpy as np
import pandas as pd
import sklearn
import xgboost as xgb
import matplotlib
```

```
In [79]: #Download the Dataset and store in /data/raw
url = 'https://github.com/gastonstat/CreditScoring/raw/master/CreditScoring.csv'
save_file_path = os.path.join('data', 'raw', 'CreditScoring.csv')

#Download the file if it does not already exist
if not os.path.exists(save_file_path):
    file_stream = requests.get(url, allow_redirects=True, stream=True)
    with open(save_file_path, 'wb+') as save_file:
        #save_file.write(file_stream)
        for chunk in file_stream.iter_content(chunk_size=1024 * 8):
            if chunk:
                save_file.write(chunk)
                save_file.flush()
                os.fsync(save_file.fileno())
```

```
In [80]: #import the csv file into a pandas dataframe
column_names = ["Status", "Seniority", "Home", "Time", "Age", "Marital", "Records", "Job", "Expenses", "Income", "Assets", "Debt", "Amount", "Price"]

df = pd.read_csv(save_file_path, names=None)
df.head()
```

Out[80]:

	Status	Seniority	Home	Time	Age	Marital	Records	Job	Expenses	Income	Assets	Debt
0	1	9	1	60	30	2	1	3	73	129	0	0
1	1	17	1	60	58	3	1	1	48	131	0	0
2	2	10	2	36	46	2	2	3	90	200	3000	0
3	1	0	1	60	24	1	1	1	63	182	2500	0
4	1	0	1	36	26	1	1	1	46	107	0	0

```
In [81]: df.describe()
```

Out[81]:

	Status	Seniority	Home	Time	Age	Marital	Records
count	4455.000000	4455.000000	4455.000000	4455.000000	4455.000000	4455.000000	4455.000000
mean	1.281257	7.987205	2.657015	46.441751	37.077666	1.879012	1.173513
std	0.450162	8.173444	1.610467	14.655225	10.984856	0.643748	0.378733
min	0.000000	0.000000	0.000000	6.000000	18.000000	0.000000	1.000000
25%	1.000000	2.000000	2.000000	36.000000	28.000000	2.000000	1.000000
50%	1.000000	5.000000	2.000000	48.000000	36.000000	2.000000	1.000000
75%	2.000000	12.000000	4.000000	60.000000	45.000000	2.000000	1.000000
max	2.000000	48.000000	6.000000	72.000000	68.000000	5.000000	2.000000

```
In [82]: #Check how many nulls there are, and decide a manner of filling these nulls in.
df.isna().sum()
```

```
Out[82]: Status      0
Seniority    0
Home         0
Time         0
Age          0
Marital      0
Records      0
Job          0
Expenses     0
Income       0
Assets       0
Debt         0
Amount       0
Price        0
dtype: int64
```

```
In [83]: #reduce the column names to lower case for convention's sake
df.columns = df.columns.str.lower()
df.head()
```

```
Out[83]:
```

	status	seniority	home	time	age	marital	records	job	expenses	income	assets	debt	amount
0	1	9	1	60	30	2	1	3	73	129	0	0	
1	1	17	1	60	58	3	1	1	48	131	0	0	
2	2	10	2	36	46	2	2	3	90	200	3000	0	2
3	1	0	1	60	24	1	1	1	63	182	2500	0	
4	1	0	1	36	26	1	1	1	46	107	0	0	

The column significance, according to the dataset source:

1 Status --- credit status 2 Seniority --- job seniority (years) 3 Home --- type of home ownership 4 Time --- time of requested loan 5 Age --- client's age 6 Marital --- marital status 7 Records --- existence of records 8 Job --- type of job 9 Expenses --- amount of expenses 10 Income --- amount of income 11 Assets --- amount of assets 12 Debt --- amount of debt 13 Amount --- amount requested of loan 14 Price --- price of good

Status is a categorical field that has been encoded as follows: "1" means "OK", and the value "2" means "default", and "0" means that the value is missing.

In [84]: *#make dictionaries that will map from the encoding to their categorical values to*

```
#for column status
status_column_mapping = {
    0 : 'unknown',
    1 : 'ok',
    2 : 'default'
}

#for column home
home_column_mapping = {
    1: 'rent',
    2: 'owner',
    3: 'private',
    4: 'ignore',
    5: 'parents',
    6: 'other',
    0: 'unknown'
}

#for column marital
marital_column_mapping = {
    1: 'single',
    2: 'married',
    3: 'widow',
    4: 'separated',
    5: 'divorced',
    0: 'unknown'
}

#for column records
records_column_mapping = {
    1: 'no',
    2: 'yes',
    0: 'unknown'
}

#for column job
job_column_mapping = {
    1: 'fixed',
    2: 'parttime',
    3: 'freelance',
    4: 'others',
    0: 'unknown'
}

def dictForStr(dict_name: str) :
    if dict_name=='home':
        return home_column_mapping
    elif dict_name=='job':
        return job_column_mapping
    elif dict_name=='status':
        return status_column_mapping
    elif dict_name=='marital':
        return marital_column_mapping
    elif dict_name=='records':
```

```

        return records_column_mapping

df_full = df.copy()
cols = ['home', 'job', 'status', 'marital', 'records']
for c in cols:
    df_full[c] = df_full[c].map(dictForStr(c))

df_full.head()

```

Out[84]:

	status	seniority	home	time	age	marital	records	job	expenses	income	assets	debt
0	ok	9	rent	60	30	married	no	freelance	73	129	0	C
1	ok	17	rent	60	58	widow	no	fixed	48	131	0	C
2	default	10	owner	36	46	married	yes	freelance	90	200	3000	C
3	ok	0	rent	60	24	single	no	fixed	63	182	2500	C
4	ok	0	rent	36	26	single	no	fixed	46	107	0	C

```
In [85]: for col in ['home', 'job', 'status', 'marital', 'records'] :  
         print("value_count for column", col, "\n",df_full[col].value_counts(), "\n---")
```

```
value_count for column home  
owner      2107  
rent       973  
parents    783  
other      319  
private    247  
ignore     20  
unknown     6  
Name: home, dtype: int64  
-----
```

```
value_count for column job  
fixed      2806  
freelance  1024  
parttime   452  
others     171  
unknown     2  
Name: job, dtype: int64  
-----
```

```
value_count for column status  
ok         3200  
default    1254  
unknown     1  
Name: status, dtype: int64  
-----
```

```
value_count for column marital  
married    3241  
single     978  
separated  130  
widow       67  
divorced    38  
unknown     1  
Name: marital, dtype: int64  
-----
```

```
value_count for column records  
no         3682  
yes        773  
Name: records, dtype: int64  
-----
```

```
In [86]: #status is the label, hence unknown values are removed
df_full = df_full[df_full.status != 'unknown']

#ensure the label col has usable values
df_full['status'].value_counts()
```

```
Out[86]: ok          3200
         default    1254
         Name: status, dtype: int64
```

```
In [87]: #columns income, assets and debt are numeric columns that have the value 99999999
#convert these to np.nan for a clearer understanding
for c in ['income', 'assets', 'debt']:
    df_full[c] = df_full[c].replace(to_replace=99999999, value=np.nan)
```

```
In [88]: #View the distribution of the numerical columns of the full, un-encoded dataset
df_full.describe().round()
```

```
Out[88]:
```

	seniority	time	age	expenses	income	assets	debt	amount	price
count	4454.0	4454.0	4454.0	4454.0	4420.0	4407.0	4436.0	4454.0	4454.0
mean	8.0	46.0	37.0	56.0	131.0	5404.0	343.0	1039.0	1463.0
std	8.0	15.0	11.0	20.0	86.0	11574.0	1246.0	475.0	628.0
min	0.0	6.0	18.0	35.0	0.0	0.0	0.0	100.0	105.0
25%	2.0	36.0	28.0	35.0	80.0	0.0	0.0	700.0	1117.0
50%	5.0	48.0	36.0	51.0	120.0	3000.0	0.0	1000.0	1400.0
75%	12.0	60.0	45.0	72.0	165.0	6000.0	0.0	1300.0	1692.0
max	48.0	72.0	68.0	180.0	959.0	300000.0	30000.0	5000.0	11140.0

At this point, since the dataset is well-understood and labels are ok, splitting the dataset between test and training sets will be done. After which cleaning and additional data-prep for training will follow...

Split Strategy:: Training -> 60%, Validation -> 20%, testing -> 20%

In [89]: *#use train_test_split from scikit-learn to split the data. First split between 60% and then the 40% block will be split in two 20% blocks.*

```
from sklearn.model_selection import train_test_split

df_train_, df_test = train_test_split(df_full, train_size=0.8, test_size=0.2)
df_train, df_validation = train_test_split(df_train_, train_size=0.75, test_size=0.25)

print(f"The num of records in test set is {len(df_test)}, while the training set  
and the validation set has {len(df_validation)} records")
```

The num of records in test set is 891, while the training set has 2672 records and the validation set has 891 records

status == 'default' when the loan was defaulted on. The model is meant to predict loan defaults. Therefore the label column will have a numeric value of 1 when there is a default (i.e., status == 1, when status == 'default' and status == 0 otherwise).


```
In [90]: y_train_ = df_train[df_train['status']=='default']
y_val_ = df_validation[df_validation['status']=='default']

y_train = (df_train.status == 'default').values
y_val = (df_validation.status == 'default').values

y_train
print(y_val)
```

```
[False False True True False False False True True False True True
 False False False False False False True True False False False True
 True False False True False False False False False False False False
 False False False False False False True True True False False True
 False False False True False True False True True False True True
 False False True False True False True False True True True True
 False False False False False True True True False True False False
 True False False False False True True True True False True False
 True True False False False False False False True False True True
 False True False False True False True False True True True False
 False False False False False False False False False False False
 False True False False True False True False True True False False
 True True True True False True False True True False False False
 False False True False False True False False True True False False
 True False False False False True False True False False False False
 False False False False False True False False False False True False
 False True True False False True False False False False True False
 False False False False True False True True True False False False
 True False True True False True False False False False True False
 False True False False True False True True True True False False
 False False True False False False False False False False True
 True True False True False False False False True True True False
 True False False True False False False False False False False False
 True False False False False False False True True False False True
 False False False False False False False False False False False
 True True False False False False True True False True False False
 False True False False False False True False False True False True
 False False False False False False True True True False False
 False False False False False True False False False True False False
```

```

False False False True False False False False False True False True
True True False True True True True False False False False False
True False True False False True False False False True True False
False True True True False False False False True True True False
True True True True False False True False False False True False
False True True False False False False False False True False False
False True False False False False True False False True False False
False False False False False False False True False False False False
True False False True True True True False True False False True
False False False False False False False False False False False False
False False False False False False False False False True True False
False False True False False True True True True True False True
True False False True False False True False False False False False
False False False True True False True False False True False False
True False False False True False True False False True False False
True False False False False False False False False False False True
False True False True False True False False False False False True
False False False False False True True True False True False False
False False False False True True False False False False False False
False False False False False True True False False True False False
False True False False True False False False False False True False
False False True True False False False False False False False False
False False True True False False False False False False True False
True True False True False True False False False False False False
False False False False False False False False True False True False
False False False True False False False False False False False False
False False False False True False False False False False False False
False True False]

```

```

In [91]: #drop the status column to ensure the model does not end up training on this column
del df_train['status']
del df_validation['status']

df_train.head()

```

Out[91]:

	seniority	home	time	age	marital	records	job	expenses	income	assets	de
2778	10	parents	36	37	single	yes	fixed	35	150.0	0.0	0
3381	7	other	60	27	single	no	fixed	35	78.0	0.0	0
46	1	owner	60	49	widow	no	freelance	35	233.0	15900.0	4500
1008	3	owner	24	29	married	no	fixed	45	254.0	6000.0	800
504	5	other	24	41	separated	no	freelance	35	0.0	0.0	0

```
In [92]: #deal with missing values: fill all NaNs with 0
#the alternative would be to plug in various kinds of central-placement figures like the mean
#I'm choosing to make this 0 instead of placing computed values that don't actually make sense

df_train = df_train.fillna(0)
df_validation = df_validation.fillna(0)

#check how many na-s we have among the numeric columns
num_of_na = 0
for col in ['seniority', 'time', 'expenses', 'income', 'assets', 'debt', 'amount']:
    num_of_na += df_train[col].isna().sum()

print(f"The total number of NaN in the numeric columns is {num_of_na}")
```

The total number of NaN in the numeric columns is 0

```
In [93]: #turn each row into a dictionary
dict_train = df_train.to_dict(orient='records')
dict_val = df_validation.to_dict(orient='records')

#print out one row/dict to see structure
print(dict_train[0])

{'seniority': 10, 'home': 'parents', 'time': 36, 'age': 37, 'marital': 'single', 'records': 'yes', 'job': 'fixed', 'expenses': 35, 'income': 150.0, 'assets': 0.0, 'debt': 0.0, 'amount': 1000, 'price': 2484}
```

```

In [94]: #The above created dictionaries can be fed into sklearn's DictVectorizer
#from sklearn documentation:
# This transformer turns lists of mappings (dict-like objects) of feature names
# scipy.sparse matrices for use with scikit-learn estimators.
# When feature values are strings, this transformer will do a binary one-hot (aka
# is constructed for each of the possible string values that the feature can take
# Note that this transformer will only do a binary one-hot encoding when feature
# features are represented as numeric values such as int or iterables of strings
# OneHotEncoder to complete binary one-hot encoding.

from sklearn.feature_extraction import DictVectorizer

#create DictVectorizer object
dv = DictVectorizer(sparse=False) #setting sparse to true would produce scipy.sparse

#make feature
X_Train = dv.fit_transform(dict_train)
X_val    = dv.transform(dict_val)

print(X_val)

[[3.1e+01 6.0e+02 4.5e+03 ... 1.0e+00 1.2e+01 1.8e+01]
 [3.0e+01 5.0e+02 7.0e+03 ... 1.0e+00 7.0e+00 3.0e+01]
 [1.9e+01 1.3e+03 9.5e+03 ... 0.0e+00 1.0e+00 6.0e+01]
 ...
 [6.3e+01 7.0e+02 3.5e+03 ... 0.0e+00 3.7e+01 4.8e+01]
 [3.2e+01 1.4e+03 5.5e+03 ... 0.0e+00 2.0e+00 6.0e+01]
 [4.3e+01 1.1e+03 3.0e+03 ... 1.0e+00 2.0e+00 6.0e+01]]

#Decision Tree Section

```

```

In [95]: #Since this is a classification task the exact Decision Tree algo being used is t
from sklearn.tree import DecisionTreeClassifier

#create the DecisionTreeClassifier object and fit the data
dt = DecisionTreeClassifier()
dt.fit(X_Train, y_train)

```

```

Out[95]: DecisionTreeClassifier()

```

```
In [96]: #Now that the model has been trained, it's performance is to be evaluated
#AUC (area under the ROC Curve) is being used as it's a well-regarded tool for bi
from sklearn.metrics import roc_auc_score

#obtain scores to evaluate with auc
y_pred = dt.predict_proba(X_Train)[:,-1]
pred_on_training = roc_auc_score(y_train, y_pred)

y_pred = dt.predict_proba(X_val)[:,-1]
pred_on_validation = roc_auc_score(y_val, y_pred)

print(f"The score against training set is {pred_on_training*100:.4f}%",
      f"while the score against validation set is {pred_on_validation*100:.4f}%."
```

The score against training set is 100.0000%, while the score against validation set is 64.6655%.

Observation: The great efficiency of the model on the training set, as opposed to the validation set, seems to show overfitting. Thus there seems to be a relative lack of the ability to generalize to unknown data.

```
In [97]: #tweak the max_depth parameter of the DecisiosnTreeClassifier to reduce over-fitt
#max_depth controls the complexity of the tree by putting a limit to the number o
#Lesser complexity is hoped to prevent over-fitting of the model to the training

dt = DecisionTreeClassifier(max_depth=2)
dt.fit(X_Train, y_train)
```

Out[97]: DecisionTreeClassifier(max_depth=2)

```
In [98]: #Since the tree itself is being controlled, it makes sense to see the tree generc

from sklearn.tree import export_text

tree_as_text = export_text(dt, feature_names=dv.feature_names_)
print(tree_as_text)
```

```
|--- seniority <= 2.50
|   |--- records=no <= 0.50
|   |   |--- class: True
|   |--- records=no > 0.50
|   |   |--- class: False
|--- seniority > 2.50
|   |--- records=no <= 0.50
|   |   |--- class: False
|   |--- records=no > 0.50
|   |   |--- class: False
```

```
In [99]: #obtain scores to evaluate with auc
y_pred = dt.predict_proba(X_Train)[: ,1]
pred_on_training = roc_auc_score(y_train, y_pred)

y_pred = dt.predict_proba(X_val)[: , 1]
pred_on_validation = roc_auc_score(y_val, y_pred)

print(f"The score against training set is {pred_on_training*100:.4f}%",
      f"while the score against validation set is {pred_on_validation*100:.4f}%."
```

The score against training set is 71.4412%, while the score against validation set is 72.1128%.

The lessened score against the training set and the increased score against the validation set indicates lessening of overfitting due to the max_depth parameter being set to 2.

In [100]: *#To continue the process, the model will be tuned for parameter according to two*

```
print('Tune Max_Depth')
#tune max_depth
for depth in list(range(1, 10, 2)):
    dt = DecisionTreeClassifier(max_depth=depth)
    dt.fit(X_Train, y_train)
    y_pred = dt.predict_proba(X_val)[:, 1]
    auc = roc_auc_score(y_val, y_pred)
    print(f"for depth of {depth}, auc is {auc*100:.4f}%")

print('\n\nTune Min_Leaf')
#tune for min_leaf_size
for depth in list(range(1, 10, 2)):
    for min_leaf_size in list(range(1, 10, 2))+[20, 30, 40, 50]:
        dt = DecisionTreeClassifier(max_depth=depth, min_samples_leaf=min_leaf_size)
        dt.fit(X_Train, y_train)
        y_pred = dt.predict_proba(X_val)[:, 1]
        auc = roc_auc_score(y_val, y_pred)
        print(f"for depth of {depth} and min_samples_leaf of {min_leaf_size}, the auc is {auc*100:.4f}%")
```

Tune Max_Depth

```
for depth of 1, auc is 64.3967%
for depth of 3, auc is 74.6159%
for depth of 5, auc is 77.5006%
for depth of 7, auc is 71.6741%
for depth of 9, auc is 66.6904%
```

Tune Min_Leaf

```
for depth of 1 and min_samples_leaf of 1, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 3, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 5, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 7, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 9, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 20, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 30, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 40, the auc is 64.3967%
for depth of 1 and min_samples_leaf of 50, the auc is 64.3967%
for depth of 3 and min_samples_leaf of 1, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 3, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 5, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 7, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 9, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 20, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 30, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 40, the auc is 74.6159%
for depth of 3 and min_samples_leaf of 50, the auc is 74.6159%
for depth of 5 and min_samples_leaf of 1, the auc is 77.4406%
for depth of 5 and min_samples_leaf of 3, the auc is 76.8247%
for depth of 5 and min_samples_leaf of 5, the auc is 77.0185%
for depth of 5 and min_samples_leaf of 7, the auc is 77.0003%
for depth of 5 and min_samples_leaf of 9, the auc is 77.0003%
for depth of 5 and min_samples_leaf of 20, the auc is 77.5650%
```

```

for depth of 5 and min_samples_leaf of 30, the auc is 76.9824%
for depth of 5 and min_samples_leaf of 40, the auc is 78.2054%
for depth of 5 and min_samples_leaf of 50, the auc is 78.6493%
for depth of 7 and min_samples_leaf of 1, the auc is 72.0631%
for depth of 7 and min_samples_leaf of 3, the auc is 74.7873%
for depth of 7 and min_samples_leaf of 5, the auc is 76.7503%
for depth of 7 and min_samples_leaf of 7, the auc is 76.6881%
for depth of 7 and min_samples_leaf of 9, the auc is 76.9175%
for depth of 7 and min_samples_leaf of 20, the auc is 78.6645%
for depth of 7 and min_samples_leaf of 30, the auc is 78.0176%
for depth of 7 and min_samples_leaf of 40, the auc is 79.2406%
for depth of 7 and min_samples_leaf of 50, the auc is 80.1572%
for depth of 9 and min_samples_leaf of 1, the auc is 67.4410%
for depth of 9 and min_samples_leaf of 3, the auc is 70.6180%
for depth of 9 and min_samples_leaf of 5, the auc is 74.3395%
for depth of 9 and min_samples_leaf of 7, the auc is 75.8207%
for depth of 9 and min_samples_leaf of 9, the auc is 75.9763%
for depth of 9 and min_samples_leaf of 20, the auc is 78.0437%
for depth of 9 and min_samples_leaf of 30, the auc is 78.5501%
for depth of 9 and min_samples_leaf of 40, the auc is 79.2610%
for depth of 9 and min_samples_leaf of 50, the auc is 80.5574%

```

```

In [101]: #The best auc is given by the following combination
#for depth of 9 and min_samples_leaf of 20, the auc is 80.4157%

#train final model with those parameters
dt = DecisionTreeClassifier(max_depth=9, min_samples_leaf=20)
dt.fit(X_Train, y_train)

```

```

Out[101]: DecisionTreeClassifier(max_depth=9, min_samples_leaf=20)

```

Random Forest Section follows...

Using the concept of Ensemble Learning, a number of models will be used to derive the verdict. The majority decision of all the verdicts will be taken to be the decision of the ensemble system as a whole.

Random Forests, in particular, work by implementing a number of trees each working on a unique combination of features to reach their classification decision. The majority of these then is taken to be the verdict outputted by the system as a whole. One strategy of selecting features is simply to randomly choose a set of features per tree in the forest.

```

In [102]: from sklearn.ensemble import RandomForestClassifier

#create a RandomForest object with 10 trees
#the parameter n_estimators sets the number of trees in the random-forest
rf = RandomForestClassifier(n_estimators=10, random_state=42)
rf.fit(X_Train, y_train)

```

```

Out[102]: RandomForestClassifier(n_estimators=10, random_state=42)

```



```
In [103]: #Check the performance of the forest
y_pred = rf.predict_proba(X_val)[: , 1]
print(f"The performance of the random forest against the validation set is {roc_auc}")
```

The performance of the random forest against the validation set is 78.0516%.

```
In [104]: #tune the n_estimator parameter
plot_dict = {}
for n in range(10, 201, 20):
    rf = RandomForestClassifier(n_estimators=n, random_state=42)
    rf.fit(X_train, y_train)
    y_pred = rf.predict_proba(X_val)[: , 1]
    s = roc_auc_score(y_val, y_pred)*100
    plot_dict[n] = s
    print(f"for n_estimator={n}, auc is {s:.4f}%.")
print(plot_dict)
```

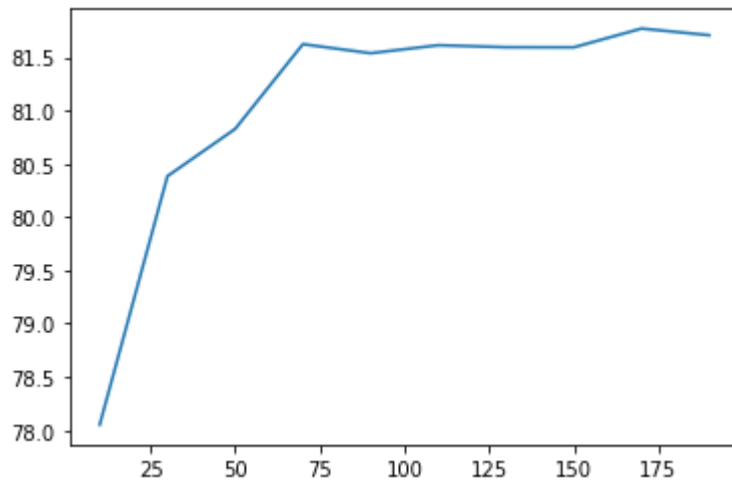
```
for n_estimator=10, auc is 78.0516%.
for n_estimator=30, auc is 80.3842%.
for n_estimator=50, auc is 80.8277%.
for n_estimator=70, auc is 81.6233%.
for n_estimator=90, auc is 81.5383%.
for n_estimator=110, auc is 81.6132%.
for n_estimator=130, auc is 81.5947%.
for n_estimator=150, auc is 81.5935%.
for n_estimator=170, auc is 81.7698%.
for n_estimator=190, auc is 81.7073%.
{10: 78.05161470406195, 30: 80.38416728358354, 50: 80.82774062792025, 70: 81.62325999101931, 90: 81.53830750373184, 110: 81.613247733589, 130: 81.59474022742995, 150: 81.59352662046871, 170: 81.7698030315902, 190: 81.70730227308584}
```

```
In [105]: #From the above the best parameter choise is given by
# ...for n_estimator=170, auc is 81.4288%.
import matplotlib.pyplot as plt

#plot: plot_dict

#plt.bar(range(len(plot_dict)), list(plot_dict.values()), align='center')
#plt.xticks(range(len(plot_dict)), list(plot_dict.keys()))
plt.plot(plot_dict.keys(), plot_dict.values())

plt.show()
```

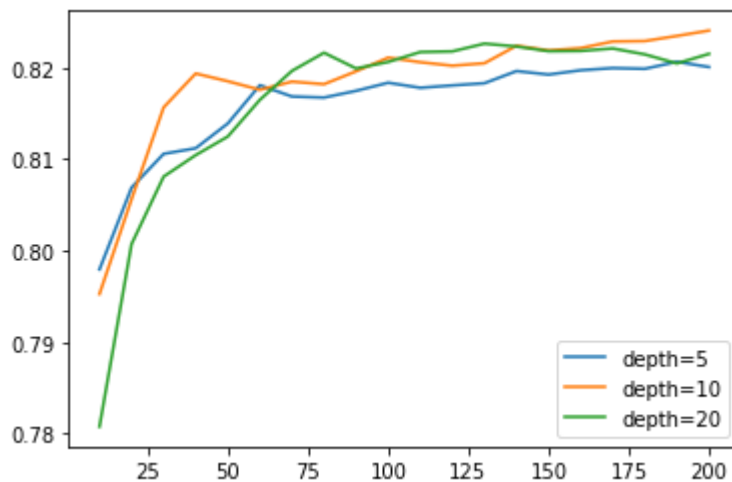


In [106]: *#tune the parameters for the trees in the random forest , i.e., max_depth and min*

```
#tune max_depth
complete_auc = {}
for depth in [5, 10, 20]:
    aucs_per_depth = []

    for i in range(10, 201, 10):
        rf = RandomForestClassifier(n_estimators=i, max_depth=depth, random_state=42)
        rf.fit(X_Train, y_train)
        y_pred = rf.predict_proba(X_val)[ :, 1]
        auc = roc_auc_score(y_val, y_pred)
        #print('%s -> %.3f' % (i, auc))
        aucs_per_depth.append(auc)
    complete_auc[depth] = aucs_per_depth

#plot auc against n_estimators per max_depth
num_trees = list(range(10, 201, 10))
plt.plot(num_trees, complete_auc[5], label='depth=5')
plt.plot(num_trees, complete_auc[10], label='depth=10')
plt.plot(num_trees, complete_auc[20], label='depth=20')
plt.legend()
plt.show()
```



max_depth of 10 shows the best performance, so lets settle on that.

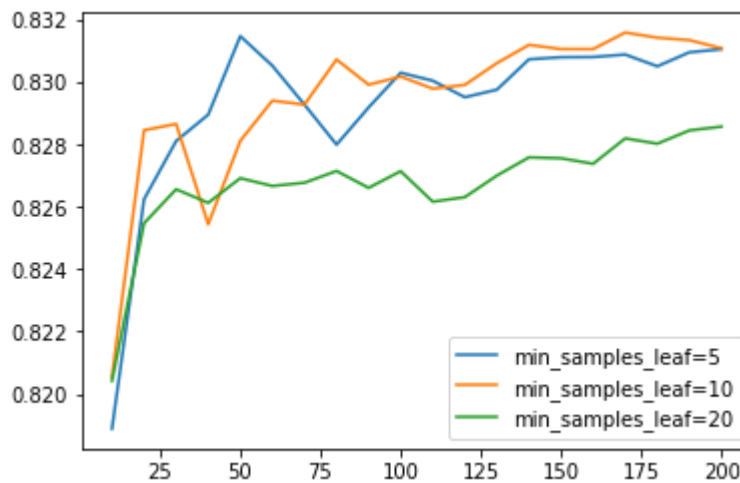
```

In [107]: #tune min_leaf_size
complete_auc = {}
for min_leaf in [5, 10, 20]:
    aucs_per_min_leaf = []

    for i in range(10, 201, 10):
        rf = RandomForestClassifier(n_estimators=i, max_depth=10, min_samples_leaf=min_leaf)
        rf.fit(X_Train, y_train)
        y_pred = rf.predict_proba(X_val)[ :, 1]
        auc = roc_auc_score(y_val, y_pred)
        #print('%s -> %.3f' % (i, auc))
        aucs_per_min_leaf.append(auc)
    complete_auc[min_leaf] = aucs_per_min_leaf

#plot auc against n_estimators per max_depth
num_trees = list(range(10, 201, 10))
plt.plot(num_trees, complete_auc[5], label='min_samples_leaf=5')
plt.plot(num_trees, complete_auc[10], label='min_samples_leaf=10')
plt.plot(num_trees, complete_auc[20], label='min_samples_leaf=20')
plt.legend()
plt.show()

```



```
In [108]: #min_samples_leaf shows best results at the value of 5
#thus, lets train the randomforest using the tuned values

rf = RandomForestClassifier(n_estimators=170, max_depth=10, min_samples_leaf=5, r
```

Gradient Boosting Section....using Extreme Gradient Boosting (xgboost) library.

Boosting refers to sequential training of models, where each model seeks to correct the deficiencies in the model that operated before it. Gradient Boosting is meant to be used with trees with particular efficiency. This differs from RandomForests in that, with random forests the individual models operate in parallel and a tally of their outputs are taken to determine the final outcome of the model.

```
In [116]: #xgboost relies on data being loaded into a DMatrix structure. This structure is
import xgboost as xgb

#create the DMatrix object for the training data
dtrain = xgb.DMatrix(X_Train, label=y_train, feature_names=dv.feature_names_)
#create the DMatrix object for the validation data
dval = xgb.DMatrix(X_val, label=y_val, feature_names=dv.feature_names_)
```

```
In [128]: #set training parameters
xgb_params = {
    'eta': 0.3, #learning rate : the weight given to correcting the output of the
    'max_depth': 6, #tree max depth parameter
    'min_child_weight': 1, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}

#set watchlist, which is a list of tuples consisting of DMatrix object and a string
#the watchlist will consist of DMatrix objects used to evaluate the model's performance
watchlist = [(dtrain, 'train'), (dval, 'val')]
```

```
In [129]: #train the actual model based on the above set parameters and data

#num_boost_round indicates the number of trees,
#set it to 10 here
model = xgb.train(xgb_params, dtrain, num_boost_round=100, evals=watchlist, verbose=0)
```

```
[14:57:07] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:541:
Parameters: { silent } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

[0]	train-auc:0.85921	val-auc:0.78128
[10]	train-auc:0.95480	val-auc:0.82279
[20]	train-auc:0.97559	val-auc:0.82299
[30]	train-auc:0.98658	val-auc:0.82289
[40]	train-auc:0.99322	val-auc:0.81959
[50]	train-auc:0.99662	val-auc:0.81963
[60]	train-auc:0.99761	val-auc:0.82005
[70]	train-auc:0.99934	val-auc:0.81690
[80]	train-auc:0.99984	val-auc:0.81633
[90]	train-auc:0.99999	val-auc:0.81766
[99]	train-auc:1.00000	val-auc:0.81554

The training set performance keeps increasing as expected with a greater number of trees. However, the performance against shows overfitting past 20 trees.

In [130]: `#tune parameters: eta, max_depth and min_child weight`

```
#tune eta:
# eta: Learning rate
#   if eta is too big, overfitting will occur soon, before the model matures well
#   if eta is too small, too many trees will need to be trained before a satisfactory model is found
#   often a value of eta=0.3 is prescribed for large datasets. This dataset is not large
# max_depth: same as seen above, the max depth the tree is allowed to take

xgb_params = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being worked on
    'max_depth': 6, #tree max depth parameter
    'min_child_weight': 1, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel training
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model = xgb.train(xgb_params, dtrain, num_boost_round=100, evals=watchlist, verbose=0)
```

[15:01:56] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:541:
Parameters: { silent } might not be used.

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

[0]	train-auc:0.85921	val-auc:0.78128
[10]	train-auc:0.91751	val-auc:0.82275
[20]	train-auc:0.94190	val-auc:0.82355
[30]	train-auc:0.95361	val-auc:0.82077
[40]	train-auc:0.96239	val-auc:0.82349
[50]	train-auc:0.97009	val-auc:0.82165
[60]	train-auc:0.97598	val-auc:0.82153
[70]	train-auc:0.98125	val-auc:0.82021
[80]	train-auc:0.98456	val-auc:0.81979
[90]	train-auc:0.98705	val-auc:0.81945
[99]	train-auc:0.98891	val-auc:0.81862

In [131]: *#slight increase in performance with the modified eta of 0.1 given by: [20] train*

#tune max_depth: to 3 and 8

min_child_weight: min num. of observations in each group

```
xgb_params1 = {  
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo  
    'max_depth': 3, #tree max depth parameter  
    'min_child_weight': 1, #min num. of observations in each group  
    'objective': 'binary:logistic', #binary classification usage  
    'eval_metric': 'auc', #the evaluation metric  
    'nthread': 8, #num. of threads used in training the model, to allow parallel  
    'seed': 42, #random state generator - set for reproducibility  
    'silent': 1  
}
```

```
model1 = xgb.train(xgb_params1, dtrain, num_boost_round=100, evals=watchlist, ver
```

```
xgb_params2 = {  
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo  
    'max_depth': 8, #tree max depth parameter  
    'min_child_weight': 1, #min num. of observations in each group  
    'objective': 'binary:logistic', #binary classification usage  
    'eval_metric': 'auc', #the evaluation metric  
    'nthread': 8, #num. of threads used in training the model, to allow parallel  
    'seed': 42, #random state generator - set for reproducibility  
    'silent': 1  
}
```

```
model2 = xgb.train(xgb_params2, dtrain, num_boost_round=100, evals=watchlist, ver
```

[15:05:59] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:

Parameters: { silent } might not be used.

This may not be accurate due to some parameters are only used in language bindings but

passed down to XGBoost core. Or some parameters are not used but slip through this

verification. Please open an issue if you find above cases.

[0]	train-auc:0.77023	val-auc:0.74305
[10]	train-auc:0.83535	val-auc:0.80944
[20]	train-auc:0.86023	val-auc:0.82033
[30]	train-auc:0.87601	val-auc:0.82754
[40]	train-auc:0.88668	val-auc:0.82910
[50]	train-auc:0.89335	val-auc:0.83202
[60]	train-auc:0.89893	val-auc:0.83166
[70]	train-auc:0.90400	val-auc:0.83382
[80]	train-auc:0.90700	val-auc:0.83534
[90]	train-auc:0.90992	val-auc:0.83551
[99]	train-auc:0.91260	val-auc:0.83480

[15:05:59] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:

Parameters: { silent } might not be used.

This may not be accurate due to some parameters are only used in language bin

things but

passed down to XGBoost core. Or some parameters are not used but slip through this

verification. Please open an issue if you find above cases.

[0]	train-auc:0.89823	val-auc:0.77731
[10]	train-auc:0.96442	val-auc:0.81075
[20]	train-auc:0.98214	val-auc:0.81542
[30]	train-auc:0.98925	val-auc:0.81375
[40]	train-auc:0.99304	val-auc:0.81527
[50]	train-auc:0.99641	val-auc:0.81479
[60]	train-auc:0.99760	val-auc:0.81515
[70]	train-auc:0.99849	val-auc:0.81566
[80]	train-auc:0.99880	val-auc:0.81725
[90]	train-auc:0.99909	val-auc:0.81775
[99]	train-auc:0.99940	val-auc:0.81690

The best performance given by eta:0.1, max_depth:3 --> [90] train-auc:0.90992 val-auc:0.83551
this likely shows benefits to limiting the tree-depth in reducing overfitting

```

In [133]: #tune min_child_weight: to 1,10 and 30
          # min_child_weight: min num. of observations in each group

xgb_params1 = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 1, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model1 = xgb.train(xgb_params1, dtrain, num_boost_round=100, evals=watchlist, ver

xgb_params2 = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 10, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model3 = xgb.train(xgb_params2, dtrain, num_boost_round=100, evals=watchlist, ver

xgb_params3 = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 30, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model3 = xgb.train(xgb_params3, dtrain, num_boost_round=100, evals=watchlist, ver

```

[15:16:19] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:541:
Parameters: { silent } might not be used.

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

[0]	train-auc:0.77023	val-auc:0.74305
[10]	train-auc:0.83535	val-auc:0.80944
[20]	train-auc:0.86023	val-auc:0.82033

```
[30]    train-auc:0.87601      val-auc:0.82754
[40]    train-auc:0.88668      val-auc:0.82910
[50]    train-auc:0.89335      val-auc:0.83202
[60]    train-auc:0.89893      val-auc:0.83166
[70]    train-auc:0.90400      val-auc:0.83382
[80]    train-auc:0.90700      val-auc:0.83534
[90]    train-auc:0.90992      val-auc:0.83551
[99]    train-auc:0.91260      val-auc:0.83480
[15:16:19] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:
Parameters: { silent } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0]      train-auc:0.76953      val-auc:0.74606
[10]     train-auc:0.83855      val-auc:0.80951
[20]     train-auc:0.85991      val-auc:0.82070
[30]     train-auc:0.87058      val-auc:0.82846
[40]     train-auc:0.87910      val-auc:0.83130
[50]     train-auc:0.88577      val-auc:0.83304
[60]     train-auc:0.89237      val-auc:0.83493
[70]     train-auc:0.89619      val-auc:0.83597
[80]     train-auc:0.89956      val-auc:0.83677
[90]     train-auc:0.90196      val-auc:0.83661
[99]     train-auc:0.90421      val-auc:0.83691
[15:16:20] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:
Parameters: { silent } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0]      train-auc:0.76905      val-auc:0.74605
[10]     train-auc:0.83204      val-auc:0.80962
[20]     train-auc:0.85303      val-auc:0.81699
[30]     train-auc:0.86301      val-auc:0.82516
[40]     train-auc:0.87069      val-auc:0.82802
[50]     train-auc:0.87491      val-auc:0.82938
[60]     train-auc:0.87915      val-auc:0.83211
[70]     train-auc:0.88186      val-auc:0.83307
[80]     train-auc:0.88404      val-auc:0.83459
[90]     train-auc:0.88648      val-auc:0.83528
[99]     train-auc:0.88819      val-auc:0.83604
```

In [140]: *#min_child_weight shows the best results*

```
xgb_params1 = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 10, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model1 = xgb.train(xgb_params1, dtrain, num_boost_round=100, evals=watchlist, ver

xgb_params2 = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 10, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model2 = xgb.train(xgb_params2, dtrain, num_boost_round=99, evals=watchlist, vert

xgb_params3 = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 10, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
model3 = xgb.train(xgb_params3, dtrain, num_boost_round=150, evals=watchlist, ver
```

[15:21:02] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:
Parameters: { silent } might not be used.

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

[0]	train-auc:0.76953	val-auc:0.74606
[10]	train-auc:0.83855	val-auc:0.80951
[20]	train-auc:0.85991	val-auc:0.82070
[30]	train-auc:0.87058	val-auc:0.82846
[40]	train-auc:0.87910	val-auc:0.83130
[50]	train-auc:0.88577	val-auc:0.83304
[60]	train-auc:0.89237	val-auc:0.83493

```
[70]    train-auc:0.89619      val-auc:0.83597
[80]    train-auc:0.89956      val-auc:0.83677
[90]    train-auc:0.90196      val-auc:0.83661
[99]    train-auc:0.90421      val-auc:0.83691
[15:21:02] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:
Parameters: { silent } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0]      train-auc:0.76953      val-auc:0.74606
[10]     train-auc:0.83855      val-auc:0.80951
[20]     train-auc:0.85991      val-auc:0.82070
[30]     train-auc:0.87058      val-auc:0.82846
[40]     train-auc:0.87910      val-auc:0.83130
[50]     train-auc:0.88577      val-auc:0.83304
[60]     train-auc:0.89237      val-auc:0.83493
[70]     train-auc:0.89619      val-auc:0.83597
[80]     train-auc:0.89956      val-auc:0.83677
[90]     train-auc:0.90196      val-auc:0.83661
[98]     train-auc:0.90405      val-auc:0.83685
[15:21:03] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:
Parameters: { silent } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0]      train-auc:0.76953      val-auc:0.74606
[10]     train-auc:0.83855      val-auc:0.80951
[20]     train-auc:0.85991      val-auc:0.82070
[30]     train-auc:0.87058      val-auc:0.82846
[40]     train-auc:0.87910      val-auc:0.83130
[50]     train-auc:0.88577      val-auc:0.83304
[60]     train-auc:0.89237      val-auc:0.83493
[70]     train-auc:0.89619      val-auc:0.83597
[80]     train-auc:0.89956      val-auc:0.83677
[90]     train-auc:0.90196      val-auc:0.83661
[100]    train-auc:0.90504      val-auc:0.83671
[110]    train-auc:0.90754      val-auc:0.83747
[120]    train-auc:0.90968      val-auc:0.83642
[130]    train-auc:0.91159      val-auc:0.83654
[140]    train-auc:0.91365      val-auc:0.83556
[149]    train-auc:0.91556      val-auc:0.83523
```

```
In [141]: #taking num. trees to be 110 as seen above for best performance of 83.747%
#the final model is below:

xgb_params_final = {
    'eta': 0.1, #learning rate : reduced here to fit the smaller dataset being wo
    'max_depth': 3, #tree max depth parameter
    'min_child_weight': 10, #min num. of observations in each group
    'objective': 'binary:logistic', #binary classification usage
    'eval_metric': 'auc', #the evaluation metric
    'nthread': 8, #num. of threads used in training the model, to allow parallel
    'seed': 42, #random state generator - set for reproducibility
    'silent': 1
}
num_trees = 110
model_final = xgb.train(xgb_params_final, dtrain, num_boost_round=num_trees)
```

```
[15:24:35] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.
0/src/learner.cc:541:
Parameters: { silent } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

In []: